

1 Introduction

There are a large number of open live webcams that are accessible over the internet. Such webcams can show interesting content from various places all around the world, but it is not feasible for a single person to browse all of the webcams to search for such interesting content. Our project aims to use Machine Learning algorithms to cluster and identify webcams of interest out of a large pool of webcams.

We scrape a list of publicly accessible, non-password-protected webcams from opentopia.com. The website provides metadata about the webcam (GPS coordinates, location name, local time) as well as a URL which can be used to retrieve a snapshot of the most recent image for a webcam. We log images from 2000 webcams with a 5-minute recording period over a duration of 1 week. The goal is to train our Machine Learning algorithms on this training data set, and then test them on the live webcams.

A large majority of webcams at any one time do not display anything interesting, and simply show a static background. An interest metric that we can define is a measure of the activity that is occurring as recorded by the webcam. In order to highlight webcams that have a high interest metric, we use a background subtraction algorithm combined with some post-processing to detect foreground objects against the background recorded by the webcam. This is detailed in section 2.

Another method that we use to explore the data set is to perform K-means clustering on the webcam images to see if the webcams in the data set can be meaningfully clustered. For example, in the data set there could be a cluster of webcams that are located outdoors, a cluster of webcams that show natural scenery, a cluster of webcams that show images of towns around the world, etc. [\[! Sean to add more stuff here about context/results behind K-means!\]](#)

2 Clustering

First, we will discuss what class of categories we are trying to achieve with our clustering. Then, we describe our method of generating features which are capable of achieving these categories. Next, we will discuss the choices we made with respect to modeling the problem and our use of clustering. Finally, we will present our results and ideas for future improvement.

Our goal is to cluster webcams in ways that people will find both interesting and useful. Therefore, it wasn't good enough to achieve clusters that are merely visually obvious (such as simple color-based clustering). Instead, we wanted to capture semantic information about each webcam and cluster based on that. To accomplish this, we decided to base our features on image descriptors. As a brief introduction, an image descriptor is a distinctive vector computed over a small region of an image. There are many algorithms that detect regions and extract descriptors in a "smart" way. Good detectors find a large number of descriptors and have a high likelihood of finding the same descriptors in a scene taken under moderately different conditions (lighting, rotation, scale). Similarly, good extractors will compute descriptors that are fairly robust to scale, lighting, and rotation. Together, this means that descriptors extracted in this way tend to be reliable indicators of complex image features, and therefore also reliable indicators of basic objects within images.

We are now tasked with converting these descriptors to features. Since each frame produces a large and varying number of descriptors, and each descriptor lives in a high-dimensional real space (for example, SIFT descriptors live in \mathbb{R}^{128}), we must find a way to compress these descriptors into a single vector. To do so, we will use a Bag-of-Features (BoF) model (akin to the Bag-of-Words model used in text). In this model, we first generate a *visual vocabulary* by clustering descriptors sampled from each webcam into a much smaller number of descriptors. Then, we can compute a feature vector for each webcam by first sampling descriptors from each webcam and then computing a histogram of the closest descriptors in the visual vocabulary. In this way, we are left with feature vectors that have the same dimension as the descriptors we have computed. Once we have these features, we can cluster the webcams into categories that roughly correspond to the presence of objects.

There are quite a few design decisions within this framework. First, we had to decide how to sample descriptors from the webcams to both generate the vocabulary and compute the feature vectors. Additionally, we had to choose which keypoint detector and descriptor extractor to use. Furthermore, we had to decide how large we wanted our vocabulary to be and how many webcam clusters we wanted. Aside from these modeling choices, we also had to make choices about the clustering algorithm itself.

To retrieve a set of descriptors from each webcam, we decided to randomly sample a fixed number of frames (uniformly) from each webcam and extract descriptors from them. We did this both for generating the visual vocabulary and for computing the feature vectors. Note that this means that larger or more-feature rich images are expected to contribute more to the vocabulary and are also expected to have feature vectors with larger norms. Although complexity is something we may be interested in clustering, image size is not. We address this later in our discussion for improvements.

We experimented with three of the most popular keypoint detectors and descriptor extractors: SIFT¹, SURF², and FREAK³. According to a lecture draft⁴, interest operators (like SIFT, SURF, and FREAK) will provide reasonably distinct descriptors for specific objects. However, dense detectors perform better for category-level tasks. Recall that we've posited that objects are what we are interested in detecting. However, we are detecting them for the purpose of categorization. Therefore, we tested the above descriptor extractors using both their corresponding keypoint detectors as well as a general dense feature detector (as implemented by OpenCV).

The choice of cluster sizes was fairly experimental. A larger vocabulary gives you more power, but also increase the risk of overfitting. Similarly, a smaller vocabulary gives you generalization but may be too weak. We eventually decided on a vocabulary of size 300. The number of webcam categories experiences similar behavior. We wish for the categories to be meaningful (implying more categories) but also easily distinguishable (implying fewer categories).

For the clustering algorithm itself, we decided to use k-means. This is perhaps the most popular clustering algorithm and was chosen due to its simplicity, reasonable runtime, and tendency to produce good results. We experimented with both random initial clusters and using k-means++ to choose initial clusters, and found that k-means++ produced significantly better clusters. **TODO(seanrafferty): Experiment with different number of runs of k-means..**

TODO(seanrafferty): Results

TODO(seanrafferty): Suggestions for improvement

3 Background Subtraction and Activity Detection

In order to calculate the interest metric of a webcam, the first step is to perform background subtraction on the stream of images coming from a webcam. The goal of this is to effectively detect activity from foreground objects (for example cars, people, animals, etc.), and to do this it is necessary to train a classifier to classify foreground objects from the background.

To classify foreground objects from the background view of a webcam, we train each pixel on a Gaussian Mixture Model (GMM) in an unsupervised learning approach. Each pixel is treated as a vector in the \mathbb{R}^3 RGB space, and is given as input to the GMM to be trained in an online learning method. The number of Gaussians is adjusted adaptively. We use the BackgroundSubtractorMOG2() function in the Python OpenCV library to perform this background subtraction algorithm. The function implements the background subtraction algorithm as detailed by Z. Zivkovic⁵.

Initial testing of the interest metric as a function of the number of classified foreground pixels revealed a flaw in this naive approach: The foreground detection algorithm suffers from abrupt lighting changes, which happen very often due to many webcams adjust their aperture dynamically in response to local lighting conditions. Such occurrences cause the classifier to erroneously classify a large number of background pixels as foreground. Thus, this interest metric results in the algorithm giving many false positives, returning images that had a change in lighting conditions without any noticeable change in the background or foreground. Thus, this initial approach did not yield acceptable results. See Figure 1 for an example.

¹D. Lowe, Distinctive Image Features from Scale-Invariant Keypoints, Canada, January 2004, <https://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>

²H. Bay, T. Tuytelaars, L. Van Gool, SURF: Speeded Up Robust Features, Zurich, 2006, <http://www.vision.ee.ethz.ch/~surf/eccv06.pdf>

³A. Alahi, R. Ortiz, P. Vandergheynst, FREAK: Fast Retina Keypoint, Switzerland, 2012, <http://infoscience.epfl.ch/record/175537/files/2069.pdf>

⁴K. Grauman, B. Leibe, Visual Recognition, 2009, http://www.cs.utexas.edu/~grau-man/courses/fall2009/papers/bag_of_visual_words.pdf

⁵Z. Zivkovic, Improved adaptive Gaussian mixture model for background subtraction, International Conference Pattern Recognition, UK, August, 2004, <http://www.zoranz.net/Publications/zivkovic2004ICPR.pdf>.

In order to handle the false positives from this naive implementation, we introduce an additional post-processing step to reduce the effect of lighting conditions on the interest metric. Since the foreground detection algorithm returns a mask where background pixels are colored black and foreground pixels are colored white, foreground objects show up as white "blobs" against the black background of the mask. The main observation (or assumption) is that the blobs of objects of interest (e.g. cars, people, animals, etc.) are small relative to the size of the image. Thus, if we can base our interest metric on some function of the blobs and threshold the blobs based on their size, we can improve the background subtraction pipeline's ability to handle changes in lighting condition.

We use the `SimpleBlobDetector()` function in the Python OpenCV library to fit contours around each blob, and threshold each blob based on a maximum percentage (nominally 20%) of the total image area. Furthermore, we now define the interest metric to be a function of the total area of blobs that occupy an image foreground. This has the effect of assigning more weight to foreground objects that are more prominent (have larger area) in the webcam image, which is what is desired. Adding this post-processing step and redefining the interest metric to be a function of the total area of blobs results in an improved webcam highlighting algorithm that is more robust to lighting changes in webcams. See Figure 2 for an example.

Even though we have made the webcam highlighting more robust against lighting changes, it is still not robust against webcams that move around, causing the background to shift constantly. The resulting background model from the background subtraction algorithm is not representative of the actual background that the webcam is pointing at, resulting in poor object detection performance. One possible way to address this would be to use the blob detection post-processing step above to check if a large portion of the image has changed, and block or refresh the background model to take into account that the background has changed.

Another improvement that can be made to the interest metric is to take into include the element of time into the calculation. Currently the webcams are highlighted based on the current webcam image that has the highest foreground blob areas, without consideration for the past history of the webcams. However, one possible way to improve the performance of the webcam highlighting might be to factor in the frequency and/or duration of the observed activity - That is, highlight a webcam only when we observe sustained activity over a pre-defined duration of time.

4 Further Work

In light of the results of our two efforts, we will begin work on labeling moving objects. Since the results of activity detection are promising, we can extract moving objects from frames by extracting a small patch containing the object. Then, we can feed these small patches to an image classifier which will label the object (e.g. "vehicle", "human", "animal", etc.). Experimenting with different classifiers will be the most challenging part of this extension. We will begin with two main approaches. First, we will use the same Bag-of-Features model that we used in clustering. Second, we will use a convolutional neural network. Since we are only interested in moving objects, our target space is greatly reduced. We do not foresee any need to detect more than five (and certainly no more than ten) different classes of objects. To train these models, we will use another labeled dataset containing objects we are interested in detecting, as well as negative examples. As a specialization to our task, we plan to focus on images that are observed by our motion detector to have a high false-positive rate (e.g. clouds, water, etc.).

We believe this extension will provide many benefits to our original goal: finding interesting and similar webcams. First, knowing the density of moving objects in each webcam enables another method of clustering which has strong semantic power. Similarly, users can search through webcams to find objects they are interested in. Additionally, having a classifier that can provide strong negative classifications (i.e. there are no interesting objects in the given patch) can help reduce the number of false positives reported by the motion detector.

5 Appendix




Figure 1: Failure of Naive background subtraction model to handle lighting changes




Figure 2: Example of background/foreground segmentation algorithm